

# Introduction to NetCDF

---

Luke Davis

Department of Atmospheric Science  
Colorado State University

# Introduction

---

Introduction

●○○○○○○○

Command-line tools

○○○○○○○○○○

Python tools

○○

Examples

○○

# Who am I?

# Background

**Problem:** Earth scientists work with huge, 4+ dimensional datasets (3 spatial dimensions, 1 time dimension, and sometimes even more).

- Example: Air temperature (longitude, latitude, height, time)
- Example: Surface pressure from “ensemble” of model simulations (longitude, latitude, time, ensemble member)
- Example: Satellite-retrieved radiation (projection x-coordinate, projection y-coordinate, wavenumber)

# Background

**Question:** What's the best way to **store** this type of data?

- Spreadsheets? Not enough dimensions.
- Matrices? No way to annotate “rows”, “columns”, etc.

**Answer:** The [NetCDF](#) format (Network Common Data Form) developed by [UCAR/Unidata](#) (right down the road!).

- Description: **Annotated** N-dimensional matrices (arrays)
- File extension: `.nc`

There are other similar data formats ([HDF](#), [GRIB](#)), and some software can work seamlessly with different formats...but **NetCDF is your new best friend**.

# Architecture

File formats have **different “version numbers”**:

- NetCDF3 (version 3) **retired in 2008**. Unidata is **currently version 4**.
- ...however version 3 still widespread (scientists are slow to change their ways...too many other things to worry about).
- Some things in NetCDF4 are impossible in NetCDF3 (multiple unlimited dimensions, ...).
- Weird read/write bugs are sometimes due to version incompatibilities.

NetCDF3 still everywhere, so you may need to use it.

# Architecture

NetCDF files have the following features:

- Global attributes.
- Global dimensions.
- Named variables, each with its own dimensions and attributes.

# Architecture

```
$ ncdump -h example.nc
```

ADD EXAMPLE



# Software

There's lots of software for working with NetCDF (HDF, GRIB) files. First, the **command-line** tools:

- NCO (NetCDF Operators).

# Software

Next, the **python** tools:

- **netCDF4** (confusingly, this works with both Netcdf versions 3 and 4!). netCDF4 is for **low-level, fast** control.
- **xarray** (“new kid” on the block but *extremely* powerful). xarray is for **high-level, convenient** control.
- Python also has **Iris “cubes”**...but this is older, less widespread, falling out of favor (xarray intended as **improvement on “cubes”**).

## Command-line tools

---

# CDO Overview

Many of us know about NCO (**N**et**C**DF **O**perators) and NCL (**N**CAR **C**ommand **L**anguage).

A comparatively recent player is **CDO** (**C**limate **D**ata **O**perators).

- Written in C++ by Max Planck Institut. **Companion** to NCO – does things NCO can't.
- Install with Anaconda, no sudo necessary: `conda install -c conda-forge cdo=1.9.0`.
- Great documentation, easy-to-remember syntax: `cdo -mergetime 2017*.nc out.nc`.
- **Subcommand chaining**: `cdo -add -sqr -selname,u f.nc -sqr -selname,v f.nc KE.nc`.

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`



# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# CDO Uses

CDO has **hundreds** of subcommands. Here's a quick survey:

- **Regrid** from one arbitrary grid (e.g. rotated pole, hexagonal cells) to another; choose from a suite of algorithms: `cdo remapcon,destination_grid.txt in.nc out.nc`.
- Interpolate to different **vertical levels**: `cdo intlevel,1000,900,800,700,600,500,400,300,200 in.nc out.nc`
- **LLS trends** ignoring missing vals: `cdo regres -seldate,1980-01-01T00:00:00,2009-12-31T23:59:59 in.nc out.nc`.
- Monthly daily and seasonal **statistics**: `cdo ymonmean in.nc climate.nc`.
- **Spatial** operations: `zonmean; fldmean; mulc,101325 -vertmean`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`



# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# XArray

If you use python, **XArray**  $\gg$  **netCDF4**. XArray is to pandas as NetCDF files are to CSV files.

- Handy objects: Datasets (entire NetCDF file) and DataArrays (single variable within NetCDF file).
- Load Dataset with `d = xr.open_dataset("file.nc", engine="pynio")`.  
For GRIB files, use `xr.open_dataset("file.grb", engine="pynio")`  
(python 3 compatibility coming soon).
- Operate along dimensions: `d["variable"].mean(dim="longitude")`.
- Permute dimensions: `d.transpose(["time", "latitude", "longitude"])`.
- Slice dimensions: `d.sel(lat=slice(0,90)); d.sel(level=850,`

# NCL

There's another tool I didn't mention...the [NCAR Command Language](#) (NCL). This is one of my favorites!

- MATLAB: Everything is an array.
- Python: Everything is an object (or dictionary, depending on who you ask).
- **NCL**: Everything is a NetCDF-formatted dataset. If you're a geoscientist, this paradigm is pretty awesome.

Sadly, you should avoid using NCL for two reasons...

- NCL [is being deprecated](#) (Unidata developers now focusing on python tools).
- NCL is very slow...among the slowest tools (see [these simple benchmarks](#)).

# NCL

**Example:** Read XYZT temperature and wind data, save YZT eddy heat flux.

```
$ ncdump -h input.nc
```

```
dimensions:
```

```
time = UNLIMITED ; // (200 currently)
```

```
lev = 60 ;
```

```
lat = 36 ;
```

```
lon = 72 ;
```

```
variables:
```

```
...
```

```
double v(time, lev, lat, lon) ;
```

```
  v:long_name = "meridional wind" ;
```

```
  v:units = "m/s" ;
```

```
double t(time, lev, lat, lon) ;
```

```
  t:long_name = "temperature" ;
```

```
  t:units = "K" ;
```

# NCL

**Example:** Read XYZT temperature and wind data, save YZT eddy heat flux.

```
$ cat fluxes.ncl
f = addfile("input.nc", "r")
o = addfile("output.nc", "c")
t = f->t
v = f->v
ehf = dim_avg_n( \
    (t - conform(t, dim_avg_n(t, 2), (/0, 1/))) \
    * (v - conform(v, dim_avg_n(v, 2), (/0, 1/))), \
    3 \
)
copy_VarCoords(t(:, :, 0), ehf)
ehf@long_name = "eddy heat flux"
ehf@units = "K*m/s"
o->ehf = ehf
```

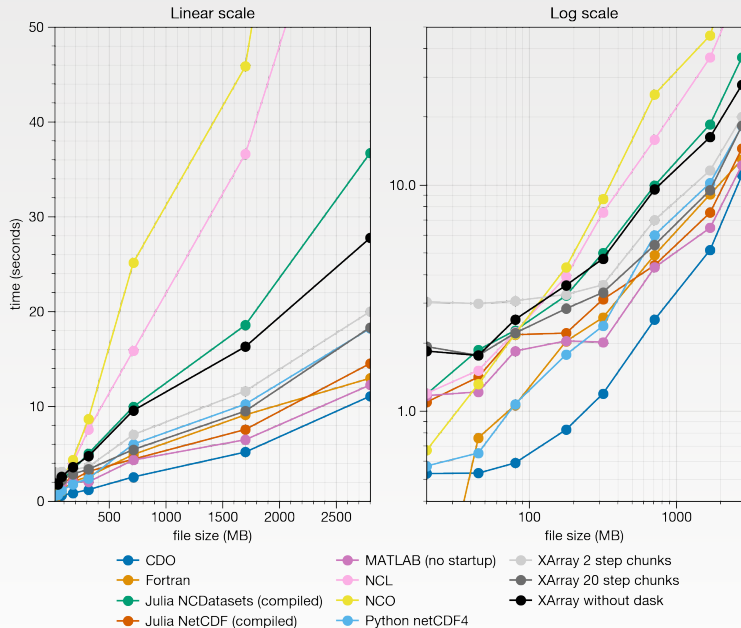
# NCL

**Example:** Read XYZT temperature and wind data, save YZT eddy heat flux.

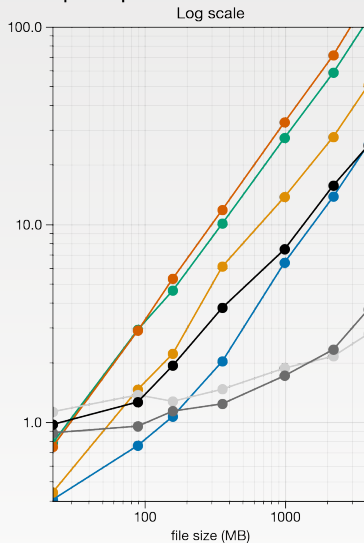

```
$ ncl fluxes.ncl
$ ncdump -h output.nc
dimensions:
  time = 200 ;
  lev = 60 ;
  lat = 36 ;
variables:
  ...
  double ehf(time, lev, lat) ;
    ehf:units = "K*m/s" ;
    ehf:long_name = "eddy heat flux" ;
```



# Fluxes benchmark on macbook

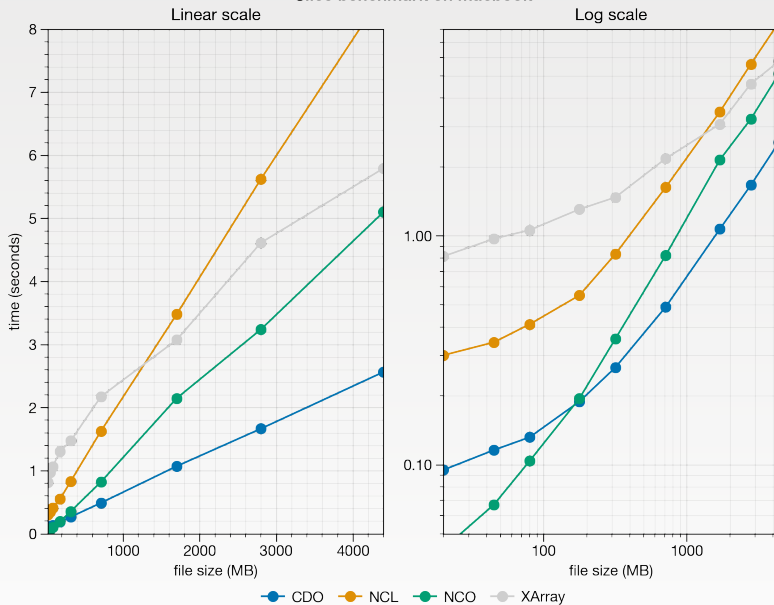


# Fluxes benchmark on supercomputer



- CDO
- Julia NCDatasets
- NCL
- NCO
- XArray 2 step chunks
- XArray 20 step chunks
- XArray without dask


# Slice benchmark on macbook



# Panoply

Many are familiar with **ncview**. Modern alternative is **Panoply**.

- Freeware released by NASA.
- Extremely easy to use.



## Python tools

---

# netCDF4

Now let's get into the python tools.

## Examples

---

## Learning more

- Use google. Try rephrasing a few times if nothing comes up.
- Use [stackoverflow](#)! ...but only *ask* questions after extensive googling fails.

Tip: For inline-code examples, use ``inline code``. For multiline-code examples, use ````multiline code````.

- Use slack! Collaboration + cooperation saves everyone time.

Tip: For inline-code and multiline-code, use backticks – just like stackoverflow.